

# RSeQC - Comprehensive Reference Guide

## Overview

RSeQC (RNA-seq Quality Control) is a comprehensive package providing quality control metrics specifically designed for aligned RNA-seq data. Unlike FastQC which operates on raw sequencing reads, RSeQC analyzes BAM alignment files to detect biology-specific issues that only become apparent after mapping to a reference genome.

**Website:** <http://rseqc.sourceforge.net/>

**Publication:** Wang et al. Bioinformatics 2012

**Current Version:** 5.0.1+

**Applicable to:** RNA-seq, small RNA-seq, long-read RNA-seq, metatranscriptomics, any BAM-based sequencing data

---

## What RSeQC Does

### Core Functions

1. **Read Distribution** - Analyzes where reads map across genomic features
2. **Junction Analysis** - Classifies splice junctions and tests saturation
3. **Gene Body Coverage** - Detects 5' to 3' bias in coverage
4. **Insert Size Distribution** - Measures fragment sizes (PE only)
5. **Artifact Profiling** - Detects systematic sequencing errors

## RSeQC vs Other QC Tools

Feature	RSeQC	FastQC	Qualimap
<b>Input</b>	BAM (aligned)	FASTQ (raw)	BAM (aligned)
<b>Splice junctions</b>	Yes	No	Limited
<b>Gene body coverage</b>	Yes	No	Yes
<b>Read distribution</b>	Detailed	No	Basic
<b>Strand detection</b>	Implicit	No	Yes
<b>Speed</b>	Moderate	Very fast	Slow
<b>RNA-seq specific</b>	Excellent	Limited	Good

---

## Installation

### Via conda (recommended)

```
conda install -c bioconda rseqc
```

## Via pip

```
pip install RSeQC
```

## Check installation

```
# List all RSeQC scripts
ls $(dirname $(which read_distribution.py))/*.py

# Check version
read_distribution.py --version
```

---

## Input Requirements

### BAM File Requirements

#### Required features:

- Sorted by coordinate
- Indexed (.bai file present)
- Contains splice junction information (CIGAR 'N' operations)
- Proper mate pairing (for PE data)

#### Verify BAM:

```
# Check if sorted
samtools view -H sample.bam | grep "S0:coordinate"

# Check for index
ls sample.bam.bai

# Check splice junctions present
samtools view sample.bam | head -100 | grep "N" | wc -l
```

### BED12 Annotation Requirements

RSeQC requires **BED12 format** (12-column BED) containing gene models:

#### Required columns:

1. Chromosome
2. Start position (0-based)
3. End position
4. Name (gene/transcript ID)
5. Score
6. Strand (+/-)
7. Thick start
8. Thick end

9. RGB color
10. Block count (exon count)
11. Block sizes (exon sizes)
12. Block starts (exon starts)

**Convert GTF to BED12:**

```
# Using UCSC tools
gtfToGenePred annotation.gtf annotation.genePred
genePredToBed annotation.genePred annotation.bed

# Using custom script (if available)
gtf2bed.py annotation.gtf > annotation.bed
```

**Verify BED12:**

```
# Check column count (must be 12)
head -n 1 annotation.bed | awk '{print NF}' 

# Check exon structure
head -n 5 annotation.bed | cut -f10-12

Critical: Chromosome names must match between BAM and BED

# Check BAM chromosomes
samtools view -H sample.bam | grep "^@SQ" | cut -f2 | sed 's/SN://'

# Check BED chromosomes
cut -f1 annotation.bed | sort -u

# If mismatch (chr1 vs 1), add/remove "chr" prefix
```

---

## RSeQC Scripts Overview

### 1. read\_distribution.py

**Purpose:** Determines where reads map across genomic features

**Command:**

```
read_distribution.py \
    --input-file sample.bam \
    --refgene annotation.bed \
    > sample.read_distribution.txt
```

**Output categories:**

Feature	Expected % (mRNA)	Interpretation
<b>CDS_Exons</b>	50-70%	Coding sequence - main target
<b>5'UTR_Exons</b>	5-10%	5' untranslated region
<b>3'UTR_Exons</b>	10-20%	3' untranslated region
<b>Introns</b>	<10%	Between exons - should be low
<b>TSS_up_1kb</b>	1-3%	Upstream of transcription start
<b>TES_down_1kb</b>	1-3%	Downstream of transcription end
<b>Intergenic</b>	<5%	Outside annotated genes

#### Example output:

Total Reads	50000000
Total Tags	52000000
Total Assigned Tags	48000000

Group	Total_bases	Tag_count	Tags/Kb
CDS_Exons	120000000	30000000	250.00
5'UTR_Exons	15000000	3000000	200.00
3'UTR_Exons	30000000	7500000	250.00
Introns	80000000	2400000	30.00
TSS_up_1kb	5000000	500000	100.00
TES_down_1kb	5000000	500000	100.00
Intergenic	100000000	2000000	20.00

#### Interpretation guide:

##### High introns (>20%):

- Genomic DNA contamination
- Degraded RNA (incomplete splicing)
- Pre-mRNA contamination
- Wrong sample type (should be RNA, not DNA)

##### High intergenic (>10%):

- Wrong genome annotation version
- Contamination with different organism
- Unannotated transcripts (could be novel - not necessarily bad)

##### Low CDS (<40%):

- Heavily degraded RNA
- Strand-specificity issues
- Wrong library type assumption

## 2. inner\_distance.py (PE only)

**Purpose:** Measures insert size distribution between paired-end mates

**Command:**

```
inner_distance.py \
    --input-file sample.bam \
    --refgene annotation.bed \
    --mapq 30 \
    --out-prefix sample
```

**Parameters:**

- **--mapq 30:** Only use high-quality alignments (99.9% confidence)
- Lower MAPQ: More reads, less reliable
- Higher MAPQ: Fewer reads, more reliable

**Formula:**

$$\text{Inner Distance} = \text{Insert Size} - (\text{Read1\_length} + \text{Read2\_length})$$

**Example:**

```
Insert size: 300bp
Read lengths: 100bp each
Inner distance: 300 - (100 + 100) = 100bp
```

**Output files:**

- `sample.inner_distance.txt`: Statistics (mean, median, std)
- `sample.inner_distance_plot.pdf`: Histogram
- `sample.inner_distance_freq.txt`: Frequency table

**Expected ranges:**

Library Type	Expected Inner Distance
mRNA (standard)	50-300bp
Small fragments	20-100bp
Long fragments	300-1000bp
Small RNA	Negative (reads overlap)

**Example output:**

```
Mean: 150bp
Median: 148bp
Std: 45bp
25th percentile: 110bp
75th percentile: 190bp
```

### Interpretation guide:

#### Very large distances (>1000bp):

- Genomic DNA contamination (spans large introns)
- Chimeric reads / wrong mate pairing
- Trans-splicing events

#### Very small distances (<20bp):

- Adapter dimers
- Over-digestion during fragmentation
- Size selection failure

#### Bimodal distribution (two peaks):

- Mixed libraries (different fragment sizes)
  - Size selection issues
  - Could indicate two distinct RNA populations (mRNA + small RNA)
- 

### 3. junction\_annotation.py

Purpose: Classifies splice junctions as known (annotated) vs novel

#### Command:

```
junction_annotation.py \
    --input-file sample.bam \
    --refgene annotation.bed \
    --mapq 30 \
    --min-intron 50 \
    --out-prefix sample
```

#### Parameters explained:

**--min-intron 50:** Minimum intron size

Gap <50bp: Deletion (sequencing error or genetic variant)

Gap 50bp: Intron (splice junction)

#### Biological rationale:

- Smallest human introns: ~50-70bp
- Typical intron: 100-10,000bp
- Largest introns: >100kb
- Setting too low: False positive junctions
- Setting too high: Miss tiny rare introns

**--mapq 30:** Mapping quality threshold

MAPQ 30 = 99.9% confidence (1 in 1000 chance of error)

MAPQ 20 = 99% confidence

MAPQ 10 = 90% confidence  
MAPQ 0 = Multi-mapper or very uncertain

Higher MAPQ:

Excludes multi-mappers (MAPQ 0-3)  
Prevents false junctions from spurious alignments  
May miss some true junctions in repetitive regions

Output files:

- sample.junction.txt: Junction statistics table
- sample.splice\_junction.pdf: Pie chart of categories
- sample.splice\_events.pdf: Splice event types
- sample.junction.bed: BED file of all detected junctions

Junction categories:

Category	Both Sites Annotated?	Interpretation
<b>Complete novel</b>	No	Could be: novel isoform, error, or wrong annotation
<b>Partial novel (donor)</b>	Donor yes, Acceptor no	Alternative acceptor site
<b>Partial novel (acceptor)</b>	Acceptor yes, Donor no	Alternative donor site
<b>Known</b>	Yes	Both splice sites in annotation

Example output:

Total junctions: 50000  
Known junctions: 42000 (84%)  
Partial novel (donor): 3000 (6%)  
Partial novel (acceptor): 3000 (6%)  
Complete novel: 2000 (4%)

Splice site motifs:

GT/AG (canonical): 47500 (95%)  
GC/AG: 2000 (4%)  
AT/AC: 300 (0.6%)  
Non-canonical: 200 (0.4%)

Interpretation guide:

>80% known junctions:

- Good quality alignment
- Correct annotation version
- Accurate junction detection

High novel junctions (>30%):

- Wrong annotation version (major update missed)
- Contamination with different organism

- Poor alignment quality (false junctions)
- Novel isoform discovery (good for research!)

**Very low junction count (<10,000 for human):**

- Genomic DNA contamination (no splicing)
- Wrong sample type
- Very low depth

**High non-canonical junctions (>5%):**

- Sequencing errors
  - Alignment artifacts
  - U12-type introns (rare but real)
- 

#### 4. junction\_saturation.py

**Purpose:** Tests if sequencing depth is sufficient to detect all junctions

**Command:**

```
junction_saturation.py \
    --input-file sample.bam \
    --refgene annotation.bed \
    --mapq 30 \
    --min-intron 50 \
    --out-prefix sample
```

**How it works:**

1. Subsample reads at 5%, 10%, 15%, ..., 95%, 100%
2. Count junctions detected at each depth
3. Plot curve of junctions vs depth
4. Assess saturation

Saturated (plateau):

More sequencing won't find many new junctions

Not saturated (still rising):

More sequencing will find more junctions

#### Output files:

- sample.junctionSaturation\_plot.pdf: Saturation curve
- sample.junctionSaturation\_plot.r: R script to recreate plot

#### Interpretation:

##### Plateau reached:

- Sufficient depth for junction detection
- Additional sequencing won't find many new junctions
- OK to proceed with analysis

##### Still rising steeply:

- Insufficient depth
- More sequencing recommended
- Proceed with caution - may miss rare isoforms

##### Never rises:

- Genomic DNA contamination (no junctions)
  - Very low quality data
- 

## 5. geneBody\_coverage.py

Purpose: Analyzes coverage uniformity across gene bodies (5' to 3')

#### Command:

```
geneBody_coverage.py \
    --input-file sample.bam \
    --refgene annotation.bed \
    --out-prefix sample
```

#### How it works:

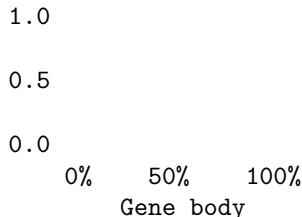
For each gene:

1. Divide into 100 bins (percentiles)
2. Calculate coverage in each bin
3. Normalize to mean coverage = 1.0
4. Average across all genes

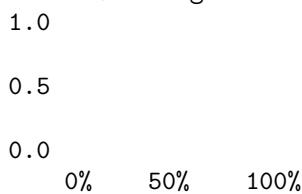
Plot shows:

X-axis: Gene position (5' → 3')  
Y-axis: Relative coverage

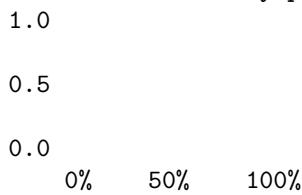
Perfect uniform coverage:



3' bias (RNA degradation):



5' bias (some library preps):



#### Output files:

- sample.geneBodyCoverage.txt: Coverage data (100 columns)
- sample.geneBodyCoverage.pdf: Coverage plot

#### Interpretation guide:

**Relatively flat profile** (coverage variation <30%):

- Good RNA quality
- Minimal degradation
- Unbiased library prep

**Strong 3' bias** (3' coverage >2x higher than 5'):

- RNA degradation
- Poly-A selection issues
- Long storage time
- Harsh RNA extraction

**Strong 5' bias** (5' coverage >2x higher than 3'):

- Some specific library prep methods
- Chemical RNA fragmentation artifacts

- Less common than 3' bias

#### **High variability:**

- Low gene count in BED file
- Annotation mismatch with BAM
- Very low sequencing depth
- Mixed sample quality

#### **Why use housekeeping genes?**

Housekeeping genes (GAPDH, ACTB, etc.):

Highly expressed → better coverage  
 Constitutively expressed → consistent across samples  
 Low variability → cleaner signal  
 Less affected by experimental conditions

All genes:

Includes low-expressed genes (noisy)  
 Includes differentially expressed genes (variable)  
 More variability in results

---

## **6. insertion\_profile.py**

**Purpose:** Analyzes where insertions occur along reads

**Command:**

```
insertion_profile.py \
    --input-file sample.bam \
    --sequencing PE \           # or SE
    --out-prefix sample
```

**What it detects:**

- Systematic insertion artifacts (not random errors)
- Position-specific insertion bias
- Library prep issues causing insertions

**Expected pattern:** Random insertions evenly distributed

**Red flags:**

- Insertions clustered at read ends → Adapter contamination
  - Insertions at specific positions → Systematic errors
  - Very high insertion rate (>2%) → Quality issues
-

## 7. deletion\_profile.py

**Purpose:** Analyzes deletion patterns along reads

**Command:**

```
deletion_profile.py \
    --input-file sample.bam \
    --out-prefix sample \
    --read-align-length 100
```

**Expected pattern:** Random deletions evenly distributed

**Red flags:**

- Deletions clustered at positions → Systematic errors
  - Very high deletion rate (>2%) → Quality issues
- 

## 8. clipping\_profile.py

**Purpose:** Analyzes where reads are soft/hard clipped

**Command:**

```
clipping_profile.py \
    --input-file sample.bam \
    --sequencing PE \
    --out-prefix sample
```

**Clipping types:**

Type	CIGAR	Meaning
Soft clip	S	Bases present in read but not aligned
Hard clip	H	Bases removed from read entirely

**Why clipping occurs:**

1. **Adapter contamination** - Clipped from ends
2. **Low quality bases** - Clipped from ends
3. **Splice junctions** - Internal clipping (RNA-seq specific)
4. **Structural variants** - Large insertions/deletions

**Expected pattern:**

- Low clipping at read ends (<5%)
- Some internal clipping (splice junctions in RNA-seq)

**Red flags:**

- High end-clipping (>20%) → Adapter issues

- Very high internal clipping → Alignment problems
- 

## 9. mismatch\_profile.py

**Purpose:** Analyzes mismatch patterns to detect systematic errors

**Command:**

```
mismatch_profile.py \
    --input-file sample.bam \
    --out-prefix sample \
    --read-align-length 100
```

**Expected pattern:** Random mismatches, low rate (<1%)

**Red flags:**

- Position-specific mismatches → Systematic sequencing errors
  - High mismatch rate (>2%) → Quality issues or wrong reference
  - Strand-specific bias → Oxidation damage (G>T on one strand)
- 

## Workflow-Specific Considerations

### RNA-seq

**Essential analyses:**

- Read distribution (detect gDNA contamination)
- Junction annotation (verify splicing)
- Gene body coverage (detect 3' bias)
- Junction saturation (assess depth)

**Expected values:**

- CDS exons: 50-70%
  - Introns: <10%
  - Known junctions: >80%
  - Uniform gene body coverage
- 

### Small RNA-seq

**Essential analyses:**

- Read distribution (should map to ncRNA features)
- Clipping profile (adapter trimming verification)

**Expected values:**

- High mapping to miRNA/snoRNA annotations
  - Negative inner distances (reads overlap)
  - Heavy adapter clipping
- 

### Long-read RNA-seq

#### Essential analyses:

- Junction annotation (long reads span multiple junctions)
- Gene body coverage (full-length transcripts)

#### Adjusted thresholds:

- Higher mismatch rate acceptable (~3-5%)
  - Different insert size expectations
  - More novel junctions (better isoform detection)
- 

### Whole Exome Sequencing (WES)

#### Essential analyses:

- Read distribution (verify on-target)
- Inner distance (check fragment sizes)

#### Expected values:

- Very high exon mapping (>90%)
- Very low intron/intergenic (<5%)
- No splice junctions (DNA, not RNA)

**Note:** Many RSeQC tools designed for RNA-seq; use with caution for DNA-seq

---

### Whole Genome Sequencing (WGS)

#### Limited applicability:

- Read distribution less meaningful (no exon enrichment)
- Junction analysis not applicable (no splicing in DNA)
- Insert size distribution useful

#### Better alternatives for WGS QC:

- Picard CollectAlignmentSummaryMetrics
  - Qualimap
  - Mosdepth
-

## Resource Requirements

### Memory (RAM)

Analysis	Full BAM (30GB)	Subsampled (1M reads)
<b>read_distribution</b>	4-8GB	2GB
<b>junction_annotation</b>	8-12GB	2GB
<b>junction_saturation</b>	12-16GB	4GB
<b>geneBody_coverage</b>	8-12GB	4GB
<b>Other profiles</b>	4-8GB	2GB

**Recommendation:** 12-16GB for full analysis

---

### CPU Cores

RSeQC scripts are **single-threaded** (no multi-core support)

#### Parallelization strategy:

- Run multiple samples in parallel
  - Run different analyses in parallel for same sample
- 

### Disk Space

#### Per sample:

- Input BAM: 5-10GB
  - Output files: 100-500MB
  - Temporary files: Minimal
- 

### Time

Analysis	Full BAM	Subsampled
<b>read_distribution</b>	5-10 min	1-2 min
<b>junction_annotation</b>	10-20 min	2-5 min
<b>junction_saturation</b>	30-60 min	5-10 min
<b>geneBody_coverage</b>	20-40 min	5-10 min
<b>All profiles</b>	10-20 min	5-10 min

**Total per sample:** ~1-2 hours (full BAM), 15-30 min (subsampled)

**For 100 samples** (parallelized): 1-3 hours

---

## Subsampling Strategy

### Why Subsample?

**Problem:** Some RSeQC analyses (junction\_saturation, geneBody\_coverage) are very slow on full BAMs

**Solution:** Subsample to 1 million reads for speed

**Benefits:**

- 5-10x faster runtime
- 90% accuracy maintained
- Still detects major issues
- Enables analysis of large cohorts

### How to Subsample

Using samtools:

```
# Calculate fraction for 1M reads
TOTAL_READS=$(samtools view -c sample.bam)
FRACTION=$(echo "scale=10; 1000000 / $TOTAL_READS" | bc)

# Subsample
samtools view -bs ${FRACTION} sample.bam > sample.1M.bam

# Index
samtools index sample.1M.bam
```

Using sambamba (faster):

```
sambamba view -f bam -t 4 --subsampling-seed=42 -s 0.05 sample.bam > sample.subsample.bam
# 0.05 = 5% = ~1M reads for 20M read dataset
```

### When to Use Full BAM

Use full BAM for:

- read\_distribution (fast anyway)
- junction\_annotation (fast anyway)
- All artifact profiles (need full data)

Use subsampled BAM for:

- junction\_saturation (very slow on full BAM)
- geneBody\_coverage (very slow on full BAM)

## **MultiQC Integration**

RSeQC outputs are automatically parsed by MultiQC for aggregated reporting.

**Files MultiQC recognizes:**

- \*.read\_distribution.txt
- \*.junction.txt
- \*.inner\_distance\_freq.txt
- \*.geneBodyCoverage.txt

**MultiQC sections created:**

- Read Distribution: Stacked bar chart (all samples)
- Gene Body Coverage: Line plot (all samples overlaid)
- Junction Annotation: Pie charts (per sample)
- Inner Distance: Histograms (per sample, PE only)

**Benefits:**

- Cross-sample comparison at a glance
  - Identify outliers easily
  - Interactive plots
  - Single HTML report
- 

## **Troubleshooting**

**"BED file format error"**

**Causes:**

- Chromosome name mismatch (BAM has "chr1", BED has "1")
- BED file not in BED12 format (wrong column count)
- Corrupted BED file

**Solutions:**

```
# Check chromosome naming
samtools view -H sample.bam | grep "^@SQ" | cut -f2 | sed 's/SN://'
cut -f1 annotation.bed | sort -u

# Verify BED12 format
head -n 1 annotation.bed | awk '{print NF}' # Should be 12

# Regenerate BED from GTF
gtfToGenePred annotation.gtf annotation.genePred
genePredToBed annotation.genePred annotation.bed
```

---

**"No reads found in region"**

**Causes:**

- BAM file empty or corrupted
- Wrong chromosome naming
- Very low mapping rate

**Solutions:**

```
# Check BAM has reads
samtools view sample.bam | head -n 10

# Check mapping rate
samtools flagstat sample.bam

# Verify chromosome names match
samtools view -H sample.bam | grep '^@SQ'"
```

---

**Script fails with "Killed" or "Out of memory"**

**Causes:**

- Very large BAM file (>50GB)
- Insufficient memory allocation
- Too many reads for analysis

**Solutions:**

- Increase memory to 16-32GB
  - Subsample BAM to 1M reads
  - Process fewer samples in parallel
- 

**"Could not determine read length"**

**Causes:**

- BAM has no aligned reads
- Alignment completely failed

**Solutions:**

```
# Check for aligned reads
samtools view sample.bam | head

# Check alignment stats
samtools flagstat sample.bam
```

---

## Gene body coverage produces weird plot

Causes:

- Too few genes in BED file
- Annotation mismatch with BAM
- Subsampling gave too few reads

Solutions:

- Check gene count: `wc -l annotation.bed`
  - Verify annotation version matches genome
  - Use more reads in subsample
- 

## Best Practices

### Standard RNA-seq QC

```
# 1. Read distribution (use full BAM)
read_distribution.py \
    --input-file sample.bam \
    --refgene annotation.bed \
    > sample.read_distribution.txt

# 2. Junction annotation (use full BAM)
junction_annotation.py \
    --input-file sample.bam \
    --refgene annotation.bed \
    --mapq 30 \
    --min-intron 50 \
    --out-prefix sample

# 3. Gene body coverage (use subsampled BAM)
geneBody_coverage.py \
    --input-file sample.1M.bam \
    --refgene housekeeping.bed \
    --out-prefix sample

# 4. Junction saturation (use subsampled BAM)
junction_saturation.py \
    --input-file sample.1M.bam \
    --refgene annotation.bed \
    --mapq 30 \
    --min-intron 50 \
    --out-prefix sample

# 5. Inner distance (PE only, use full BAM)
```

```
inner_distance.py \
--input-file sample.bam \
--refgene annotation.bed \
--mapq 30 \
--out-prefix sample
```

---

### Minimal QC (Fast)

For quick assessment:

```
# Just read distribution and junction annotation
read_distribution.py --input-file sample.bam --refgene annotation.bed > sample.txt
junction_annotation.py --input-file sample.bam --refgene annotation.bed --out-prefix sample
```

Takes ~5-10 minutes, catches most major issues.

---

### Comprehensive QC (Full)

For publication-quality analysis, run all 9 analyses as shown in the process.

---

## Related Documentation

- **FastQC:** docs/fastqc.md - Raw read QC
  - **MultiQC:** docs/multiqc.md - Aggregated reporting
  - **STAR Alignment:** docs/star\_align.md - Generating input BAMs
  - **RSeQC Manual:** <http://rseqc.sourceforge.net/>
- 

**Document Version:** 1.0

**Last Updated:** January 2026

**RSeQC Version:** 5.0.1+

**Applicable to:** All BAM-based sequencing applications, optimized for RNA-seq