

# Salmon Quantification - Comprehensive Reference Guide

## Overview

Salmon is a fast, alignment-free tool for quantifying transcript and gene expression from RNA-seq data. It uses a k-mer-based quasi-mapping approach combined with an Expectation-Maximization (EM) algorithm to estimate transcript abundances without requiring full genome alignment.

**Website:** <https://combine-lab.github.io/salmon/>

**Publication:** Patro et al. Nature Methods 2017

**Current Version:** 1.10.0+

**Applicable to:** RNA-seq, single-cell RNA-seq, metatranscriptomics, spatial transcriptomics

---

## What Salmon Does

### Core Functions

1. **Quasi-mapping** - Maps reads to transcripts using k-mer matching (not full alignment)
2. **Transcript Quantification** - Estimates abundance for each transcript (TPM and counts)
3. **Gene-level Aggregation** - Can aggregate to gene level (or use tximport)
4. **Multi-mapper Resolution** - Probabilistically assigns reads mapping to multiple transcripts
5. **Bias Correction** - Corrects for GC, sequence, and positional biases

## Salmon vs Traditional Methods

Feature	Salmon	STAR + featureCounts	kallisto
<b>Speed</b>	Very fast (2-5 min)	Slow (20-40 min)	Very fast (3-8 min)
<b>Accuracy</b>	Excellent	Excellent	Excellent
<b>Output</b>	Transcript counts	Gene counts	Transcript counts
<b>BAM file</b>	No	Yes	No
<b>Memory</b>	8-12GB	30-50GB	4-8GB
<b>Bias correction</b>	Built-in	Manual	Limited
<b>Multi-mappers</b>	EM algorithm	Discard or count once	EM algorithm
<b>Gene families</b>	Handles well	Problematic	Handles well

---

## Salmon Quantification Approach

### K-mer Based Quasi-Mapping

Traditional alignment (STAR, HISAT2):

Read → Find seeds → Extend alignment → Score → Report best alignment

Time: ~10-30 minutes per sample

Output: Full BAM file with exact positions

Salmon quasi-mapping:

Read → Hash k-mers → Find transcript matches → Score compatibility → Report

Time: ~2-5 minutes per sample

Output: Transcript abundance estimates (no BAM)

## How It Works

Step 1: K-mer Extraction

Read: ATGCGTACGATCGATCG...

K-mers (k=31): ATGCGTACGATCGATCGATCGATCGATCG

TGCGTACGATCGATCGATCGATCGATCG

GCGTACGATCGATCGATCGATCGA

...

Step 2: Transcript Matching

K-mer → Index lookup → Find transcripts containing this k-mer

Build compatibility graph: Read ↔ {Transcript1, Transcript2, ...}

Step 3: Probabilistic Assignment (EM Algorithm)

Initialize: Uniform probability across compatible transcripts

Iterate:

1. E-step: Assign reads to transcripts weighted by current estimates

2. M-step: Update transcript abundances based on assignments

Until convergence

Result: Probability distribution of each read across transcripts

→ Fractional counts for each transcript

## Why EM Algorithm?

Problem: Multi-mapping reads

Read maps to:

- Transcript A (gene X)
- Transcript B (gene X, different isoform)
- Transcript C (pseudogene of X)

Traditional methods:

```
Discard (lose information)
Count once (underestimate)
Count all (overestimate)
```

Salmon's EM:

```
Assign fractionally based on:
- Transcript abundance
- Sequence compatibility
- Fragment length distribution
Iterate until stable
Maintains total read count
```

Example:

Initial:

```
Transcript A: 1000 reads (100 unique, 900 ambiguous)
Transcript B: 100 reads (10 unique, 90 ambiguous)
```

EM realizes:

```
- A is 10x more abundant
- Ambiguous reads more likely from A
```

After EM:

```
Transcript A: ~945 reads (probabilistic assignment)
Transcript B: ~55 reads (probabilistic assignment)
```

---

## Salmon Quantification Command

### Basic Command

```
salmon quant \
  --index salmon_index \
  --libType A \
  --mates1 R1.fastq.gz \
  --mates2 R2.fastq.gz \
  --output sample_quant \
  --threads 4
```

### Command Breakdown

```
salmon quant          # Quantification mode
  --index salmon_index # Path to Salmon index
  --libType A          # Auto-detect library type
  --mates1 R1.fastq.gz # Forward reads (PE)
  --mates2 R2.fastq.gz # Reverse reads (PE)
```

```
--output sample_quant \\ # Output directory
--threads 4 # Number of CPU cores
```

**Single-end data:**

```
salmon quant \\
  --index salmon_index \\
  --libType A \\
  --unmatedReads reads.fastq.gz \\
  --output sample_quant \\
  --threads 4
```

---

## Key Parameters Explained

### Library Type (--libType)

**Purpose:** Specifies strand orientation of reads

### Auto-detection (RECOMMENDED):

```
--libType A # Auto-detect library type
```

Salmon will examine first 1-10 million reads and determine:

- **SE Libraries:** SF, SR, or U
- **PE Libraries:** ISF, ISR, or IU

### Library type codes:

Code	Description	Read Orientation
SF	Single-end Forward	Read matches transcript strand
SR	Single-end Reverse	Read is reverse complement
U	Unstranded	Read could be either orientation
ISF	Inward Stranded Forward	R2 matches transcript, R1 is reverse
ISR	Inward Stranded Reverse	R1 matches transcript, R2 is reverse
IU	Inward Unstranded	Reads could be either orientation

### Common protocols:

Protocol	Library Type
Illumina TruSeq	ISR (reverse stranded)
dUTP method	ISR (reverse stranded)
Ligation-based	ISR or ISF
SMARTer	Varies (use auto)
Nextera	Varies (use auto)
Old protocols	IU (unstranded)

### Verification:

```
# Check detected library type
cat sample_quant/lib_format_counts.json

# Example output
{
    "expected_format": "ISR",
    "compatible_fragment_ratio": 0.9823,
    "num_compatible_fragments": 9823000,
    "num_assigned_fragments": 10000000,
    ...
}
```

### If auto-detection fails:

```
# Manually specify (not recommended)
--libType ISR      # If you know it's reverse-stranded PE
```

---

## Selective Alignment (--validateMappings)

**Purpose:** More stringent mapping validation

**How it works:**

**Without --validateMappings** (default quasi-mapping):

1. Find k-mer matches
2. Count matches per transcript
3. Assign based on counts  
→ Fast but may include spurious matches

**With --validateMappings** (selective alignment):

1. Find k-mer matches
2. Build alignment chains
3. Score alignment quality
4. Filter low-quality matches
5. Assign based on validated matches  
→ ~20% slower but more accurate

**Accuracy improvement:** ~2-5% more accurate quantification

**Recommendation:**

- Always use for publication-quality results
- Can skip for quick exploratory analysis

```
salmon quant \
    --validateMappings \      # RECOMMENDED for accuracy
```

```
--index salmon_index \
...
```

---

## Bias Corrections

Salmon provides three types of bias correction:

1. **GC Bias (--gcBias) Problem:** PCR amplification favors certain GC content

Without correction:

High GC transcripts: Under-estimated  
Low GC transcripts: Over-estimated

Example:

GC-rich gene (70% GC): 500 reads → Actually 600 reads  
GC-poor gene (30% GC): 500 reads → Actually 400 reads

How Salmon corrects:

1. Model GC content vs fragment count
2. Learn bias curve from data
3. Re-weight fragments based on GC content
4. Output corrected abundances

Command:

```
--gcBias      # Enable GC bias correction
```

Impact: ~3-5% improvement in accuracy

When to use:

- PCR-amplified libraries
  - Non-UMI protocols
  - Publication-quality analysis
  - UMI-based methods (already corrected)
- 

2. **Sequence Bias (--seqBias) Problem:** Random hexamer priming isn't truly random

Random hexamer priming:

Primers: NNNNNN (supposedly random)  
Reality: Some hexamers bind better than others

Bias example:

GCGCGC: Binds strongly → Over-represented

ATATAT: Binds weakly → Under-represented

**Result:**

Transcripts starting with preferred hexamers → Over-estimated  
Transcripts starting with poor hexamers → Under-estimated

**How Salmon corrects:**

1. Learn hexamer preferences from read starts
2. Model position-specific bias
3. Correct abundances based on sequence context

**Command:**

`--seqBias # Enable sequence-specific bias correction`

**Impact:** ~2-3% improvement in accuracy

**When to use:**

- Random hexamer priming protocols
  - Any RT-PCR based method
  - Oligo-dT priming (different bias pattern)
- 

**3. Positional Bias (--posBias)** **Problem:** Fragments not uniformly distributed along transcripts

**Causes:**

- 3' bias: RNA degradation, poly-A selection
- 5' bias: Some library prep methods
- Coverage holes: Secondary structure

**Without correction:**

3' biased sample:  
Gene A (5' region expressed): Under-estimated  
Gene B (3' region expressed): Over-estimated

**How Salmon corrects:**

1. Model fragment position distribution
2. Learn coverage pattern along transcripts
3. Re-weight fragments based on position

**Command:**

`--posBias # Enable positional bias correction`

**Impact:** ~1-2% improvement (varies by sample quality)

**When to use:**

- Degraded RNA samples
  - Poly-A selected libraries (mild 3' bias)
  - High-quality, uniform coverage
- 

**Combined Bias Correction** Standard practice: Enable all three

```
salmon quant \
    --validateMappings \
    --gcBias \
    --seqBias \
    --posBias \
    --index salmon_index \
    --libType A \
    --mates1 R1.fastq.gz \
    --mates2 R2.fastq.gz \
    --output sample_quant \
    --threads 4
```

Total improvement: ~5-10% more accurate quantification

Trade-off: ~10-15% slower runtime (still very fast overall)

---

### Mapping Parameters

**--minAssignedFrags** Purpose: Minimum fragments assigned to a transcript for it to be reported

Default: 10

#### Why it matters:

Low-count transcripts (1-9 fragments):

- High uncertainty
- Likely noise or mapping artifacts
- Unstable estimates

Filtering helps:

- Reduces false positives
- Improves statistical power
- Cleaner downstream analysis

#### Recommendations:

```
# Standard analysis
--minAssignedFrags 10      # Default, good balance

# Sensitive analysis (detect low-abundance)
```

```
--minAssignedFrags 1      # Keep everything  
  
# Conservative analysis  
--minAssignedFrags 50     # Only well-supported transcripts
```

---

**--rangeFactorizationBins** **Purpose:** Controls EM algorithm precision

**Default:** 4

**What it does:**

Bins for fragment length distribution:

More bins → More precise → Slower

Fewer bins → Less precise → Faster

Default (4): Good balance

Higher (8-16): Marginal improvement, longer runtime

Lower (2): Faster but less accurate

**Recommendation:** Keep default unless you have specific needs

---

## Output Files

Salmon creates an output directory with multiple files:

### Directory Structure

```
sample_quant/  
    quant.sf                      # Main output: Transcript abundances  
    quant.genes.sf                 # Gene-level abundances (if --geneMap used)  
    lib_format_counts.json         # Library type detection results  
    meta_info.json                 # Run metadata and parameters  
    cmd_info.json                  # Command line executed  
    aux_info/  
        ambig_info.tsv            # Ambiguously mapped fragment info  
        expected_bias.gz          # Learned bias models  
        fld.gz                     # Fragment length distribution  
        meta_info.json             # Auxiliary metadata  
        observed_bias.gz          # Observed bias patterns  
    logs/  
        salmon_quant.log          # Detailed log file
```

---

## **quant.sf (Main Output)**

**Format:** Tab-separated values (TSV)

**Columns:**

Column	Name	Description
1	<b>Name</b>	Transcript ID (e.g., ENST000000000001)
2	<b>Length</b>	Transcript length in base pairs
3	<b>EffectiveLength</b>	Effective length accounting for biases
4	<b>TPM</b>	Transcripts Per Million (normalized abundance)
5	<b>NumReads</b>	Estimated number of reads from this transcript

**Example:**

Name	Length	EffectiveLength	TPM	NumReads
ENST000000000001.1	2000	1850	1250.50	10000.5
ENST000000000002.1	1500	1350	500.25	3500.2
ENST000000000003.1	3000	2850	2500.75	25000.8

**Column details:**

### **Length**

- Raw transcript length from FASTA
- Includes all exons (spliced)
- Fixed value for each transcript

### **EffectiveLength**

- Adjusted length accounting for:
  - Fragment length distribution
  - Positional bias
  - Edge effects (fragments can't start/end anywhere)

**Formula:** EffectiveLength = Length -  $\frac{1}{\text{mean fragment length}}$

**Example:**

```

Transcript length: 2000bp
Mean fragment length: 200bp
Effective length: 2000 - 200 + 1 = 1801bp

```

**Why:** A 200bp fragment can't start in last 199bp of transcript

## TPM (Transcripts Per Million)

- Normalized abundance measure
- Comparable across samples
- Sums to 1 million within each sample

Formula:

1. Counts per base: NumReads / EffectiveLength
2. Normalize: (Counts per base / Sum of all counts per base) × 1,000,000

Properties:

Sum of all TPMs = 1,000,000  
Accounts for transcript length  
Comparable across samples  
NOT raw counts (don't use for DESeq2/edgeR directly)

TPM interpretation:

TPM = 1000: 1 out of every 1000 transcripts in the sample  
TPM = 100: Moderate expression  
TPM = 10: Low expression  
TPM = 1: Very low expression  
TPM < 1: Barely detectable

## NumReads

- Estimated number of reads from transcript
- Can be fractional (due to multi-mapping)
- Sum across transcripts total mapped reads

Example:

Read maps to Transcript A and B

If A is 3x more abundant than B:  
Transcript A: +0.75 reads  
Transcript B: +0.25 reads

Result: NumReads can be non-integer

---

## lib\_format\_counts.json

Purpose: Library type detection results

Example:

```
{  
  "expected_format": "ISR",  
  "compatible_fragment_ratio": 0.9823,
```

```

    "num_compatible_fragments": 9823000,
    "num_assigned_fragments": 10000000,
    "num_consistent_mappings": 9823000,
    "num_inconsistent_mappings": 177000,
    "MSF": 50000,
    "MSR": 127000,
    "MU": 0,
    "SF": 25000,
    "SR": 25000,
    "U": 0,
    "ISF": 88500,
    "ISR": 9823000,
    "IU": 0
}

```

**Key fields:**

- `expected_format`: Detected library type
- `compatible_fragment_ratio`: Fraction of reads matching expected orientation
- High ratio (>0.95): Confident detection
- Low ratio (<0.80): Mixed or poor quality

**Red flags:**

- Compatible ratio <0.80: Library type unclear
  - Multiple types with similar counts: Mixed libraries
  - Unexpected type: Wrong protocol or sequencing issue
- 

**meta\_info.json**

**Purpose:** Run metadata and parameters

**Example:**

```
{
    "salmon_version": "1.10.0",
    "samp_type": "paired",
    "opt_type": "vbem",
    "quant_errors": [],
    "num_libraries": 1,
    "library_types": ["ISR"],
    "frag_dist_length": 1000,
    "num_targets": 234567,
    "num_bootstraps": 0,
    "mapping_type": "mapping",
    "index_seq_hash": "a1b2c3d4e5f6..."
}
```

```

    "index_name_hash": "f6e5d4c3b2a1...",
    "index_decoy_seq_hash": "1a2b3c4d5e6f...",
    "index_keep_duplicates": false,
    "keep_duplicates": false
}

```

#### Important fields:

- `salmon_version`: Version used (important for reproducibility)
  - `num_targets`: Number of transcripts quantified
  - `library_types`: Detected library type(s)
  - `num_bootstraps`: If bootstrapping was run
- 

### `aux_info/` Directory

Contains auxiliary information for downstream tools:

#### `fld.gz`

- Fragment length distribution
- Used by tximport for length correction
- Histogram of insert sizes

#### `expected_bias.gz / observed_bias.gz`

- Learned bias models
- Used for bias correction
- Can be visualized for QC

#### `ambig_info.tsv`

- Information about ambiguously mapped fragments
  - Useful for debugging low mapping rates
- 

## Downstream Analysis with tximport

### Why Use tximport?

**Problem:** Salmon outputs transcript-level abundances, but we often want gene-level

#### Solutions:

1. Sum transcript counts directly → Loses statistical information
2. Use longest isoform only → Ignores real biology
3. Use tximport → Properly aggregates with statistical awareness

## tximport Advantages

- Aggregates transcripts to genes correctly
  - Preserves inferential uncertainty
  - Accounts for transcript length differences
  - Compatible with DESeq2/edgeR
  - Handles multi-isoform genes properly
- 

## Basic tximport Usage

### Step 1: Prepare tx2gene mapping

```
library(tximport)
library(dplyr)

# Option A: From Ensembl BioMart
library(biomaRt)
mart <- useEnsembl("ensembl", dataset = "hsapiens_gene_ensembl")
tx2gene <- getBM(
  attributes = c("ensembl_transcript_id", "ensembl_gene_id"),
  mart = mart
)
colnames(tx2gene) <- c("transcript_id", "gene_id")

# Option B: From GTF file
library(GenomicFeatures)
txdb <- makeTxDbFromGFF("annotation.gtf")
k <- keys(txdb, keytype = "TXNAME")
tx2gene <- select(txdb, k, "GENEID", "TXNAME")
colnames(tx2gene) <- c("transcript_id", "gene_id")

# Option C: Manual (if you have a file)
tx2gene <- read.table("tx2gene.txt", header = TRUE)
```

Format:

```
transcript_id      gene_id
ENST000000000001  ENSG000000000001
ENST000000000002  ENSG000000000001
ENST000000000003  ENSG000000000002
```

---

### Step 2: Collect Salmon output files

```
# List all quant.sf files
samples <- c("sample1", "sample2", "sample3", ...)
```

```

files <- file.path("salmon_output", samples, "quant.sf")
names(files) <- samples

# Verify files exist
all(file.exists(files))

```

---

### Step 3: Import with tximport

```

# Import transcript abundances
txi <- tximport(
  files,
  type = "salmon",
  tx2gene = tx2gene,
  countsFromAbundance = "lengthScaledTPM" # Recommended
)

# Output structure
names(txi)
# [1] "abundance"           # Gene-level TPMs
# [2] "counts"              # Gene-level estimated counts
# [3] "length"               # Average transcript length per gene
# [4] "countsFromAbundance" # Method used

```

**countsFromAbundance** options:

Option	Description	Use Case
"no"	Use estimated counts directly	Standard
"scaledTPM"	Scale TPMs to library size	Simple normalization
"lengthScaledTPM"	Scale TPMs & adjust for length	<b>RECOMMENDED</b>
"dtuScaledTPM"	For DTU analysis	Isoform switching

---

### Step 4: Create DESeq2 object

```

library(DESeq2)

# Sample metadata
coldata <- data.frame(
  sample = samples,
  condition = c("control", "control", "treated", ...),
  row.names = samples
)

```

```

# Create DESeqDataSet from tximport
dds <- DESeqDataSetFromTximport(
  txi,
  colData = coldata,
  design = ~ condition
)

# Run differential expression
dds <- DESeq(dds)
results <- results(dds)

```

---

### Step 5: Explore results

```

# View results
head(results)

# MA plot
plotMA(results)

# PCA
vsd <- vst(dds)
plotPCA(vsd, intgroup = "condition")

# Top genes
res_ordered <- results[order(results$padj), ]
head(res_ordered, 20)

```

---

## Advanced tximport Features

### Import with bootstraps (if run)

```

# Run Salmon with bootstraps
salmon quant \
  --numBootstraps 30 \      # 30 bootstrap samples
  ...

# Import with uncertainty
txi <- tximport(
  files,
  type = "salmon",
  tx2gene = tx2gene,
  txOut = FALSE,           # Gene-level
  varReduce = TRUE          # Compute inferential variance
)

```

```
# Now txi includes variance estimates
names(txi)
# Includes "variance" element
```

#### When to use bootstraps:

- Publication-quality DEG analysis
  - When uncertainty quantification matters
  - Low-count genes
  - Exploratory analysis (overkill)
- 

#### Transcript-level analysis

```
# Keep transcript-level (don't aggregate to genes)
txi_transcripts <- tximport(
  files,
  type = "salmon",
  txOut = TRUE,           # Keep transcript-level
  countsFromAbundance = "lengthScaledTPM"
)

# Use for isoform-level analysis
dds_tx <- DESeqDataSetFromTximport(
  txi_transcripts,
  colData = coldata,
  design = ~ condition
)
```

#### Use cases:

- Isoform switching analysis
  - Alternative splicing detection
  - Transcript-specific regulation
- 

### Interpreting Salmon Logs

#### Mapping Rate

Check salmon\_quant.log:

```
[info] Mapping rate = 85.23%
```

Interpretation:

Mapping Rate	Quality	Action
>80%	Excellent	Proceed
70-80%	Good	OK, but check contamination
60-70%	Acceptable	Investigate low mapping
<60%	Poor	Check index, contamination, degradation

### Common causes of low mapping:

- Wrong index (different species or annotation version)
  - Contamination (bacterial, adapter, rRNA)
  - Degraded RNA
  - Poor sequencing quality
- 

### Fragment Length Distribution

Check aux\_info/fld.gz:

```
# Visualize in R
library(dplyr)
library(ggplot2)

fld <- read.table(gzfile("sample_quant/aux_info/fld.gz"))
colnames(fld) <- c("length", "count")

ggplot(fld, aes(x = length, y = count)) +
  geom_line() +
  labs(title = "Fragment Length Distribution",
       x = "Fragment Length (bp)",
       y = "Count")
```

Expected patterns:

Normal distribution (good):

100    200    300    (bp)

Bimodal (mixed libraries):

100    200    300

Very narrow (size selection issue):

200

Very wide (poor quality):

50      300      500

---

### Library Type Consistency

Check across samples:

```
# Extract library types from all samples
for sample in sample1 sample2 sample3; do
    echo -n "${sample}: "
    grep "expected_format" ${sample}/lib_format_counts.json
done

# Expected: All samples same type
sample1: "expected_format": "ISR"
sample2: "expected_format": "ISR"
sample3: "expected_format": "ISR"

# Red flag: Mixed types
sample1: "expected_format": "ISR"
sample2: "expected_format": "IU"    # ← Problem!
sample3: "expected_format": "ISR"
```

If inconsistent:

- Check library prep protocol
  - Verify not mixing different library types
  - May indicate sample swap or mislabeling
- 

### Resource Requirements

#### Memory (RAM)

Component	Requirement
Salmon index in RAM	5-8GB (human)
Quantification buffer	2-4GB

Component	Requirement
<b>Total</b>	8-12GB

#### Scaling by species:

- Mouse: 6-10GB
  - Zebrafish: 4-6GB
  - Drosophila: 2-4GB
  - C. elegans: 1-2GB
  - Arabidopsis: 2-3GB
- 

#### CPU Cores

#### Scaling efficiency:

- 1 core: Baseline
- 4 cores: 3.5x faster (good efficiency)
- 8 cores: 6x faster (diminishing returns)
- 16 cores: 8x faster (poor efficiency)

**Recommendation:** 4-8 cores optimal

---

#### Disk Space

#### Per sample:

- Input FASTQ: 2-8GB (compressed)
- Output directory: 50-200MB
- Temporary files: Minimal

**For 100 samples:** ~5-20GB total output

---

#### Time

Sample Size	4 cores	8 cores
<b>10M reads</b>	2 min	1.5 min
<b>25M reads</b>	4 min	2.5 min
<b>50M reads</b>	7 min	4 min
<b>100M reads</b>	12 min	7 min

**With bias correction:** +10-20% time

---

## Workflow-Specific Considerations

### Bulk RNA-seq

Standard settings:

```
salmon quant \
    --validateMappings \
    --gcBias \
    --seqBias \
    --posBias \
    --libType A \
    --index salmon_index \
    --mates1 R1.fastq.gz \
    --mates2 R2.fastq.gz \
    --threads 4 \
    --output sample_quant
```

Expected metrics:

- Mapping rate: >75%
  - Library type: Consistent across samples
  - Fragment length: 150-300bp mean
- 

### Single-cell RNA-seq

Special considerations:

```
# alevin mode (for UMI-based scRNA-seq)
salmon alevin \
    --index salmon_index \
    --libType ISR \
    --mates1 R1.fastq.gz \# Cell barcode + UMI \
    --mates2 R2.fastq.gz \# cDNA read \
    --chromiumV3 \# Or other chemistry \
    --output sample_alevin \
    --threads 8
```

Differences:

- Uses alevin mode (not quant)
  - Handles UMIs and cell barcodes
  - Outputs count matrix (genes × cells)
  - No bias correction (UMIs handle that)
-

## Metatranscriptomics

### Challenges:

- Multiple species in index
- High multi-mapping
- Taxonomic assignment needed

### Recommendations:

```
salmon quant \
    --validateMappings \
    --allowDovetail \          # For short fragments
    --recoverOrphans \         # Keep unpaired reads
    --libType A \
    --index metatranscriptome_index \
    --mates1 R1.fastq.gz \
    --mates2 R2.fastq.gz \
    --threads 8 \
    --output sample_quant
```

### Downstream:

- Use taxonomic classification
  - Aggregate by species/genus
  - Account for shared genes
- 

## Troubleshooting

### Low Mapping Rate (<50%)

#### Diagnostic steps:

1. Check index matches data:

```
# Verify species
head -n 1 salmon_index/info.json

# Check number of transcripts
grep "num_targets" sample_quant/meta_info.json
```

2. Check for contamination:

```
# Run FastQC
fastqc R1.fastq.gz R2.fastq.gz

# Check adapter content
# Check overrepresented sequences
```

3. Test with subset:

```

# Take first 100K reads
zcat R1.fastq.gz | head -400000 | gzip > R1.subset.fq.gz
zcat R2.fastq.gz | head -400000 | gzip > R2.subset.fq.gz

# Run Salmon
salmon quant --index salmon_index --libType A \
    --mates1 R1.subset.fq.gz --mates2 R2.subset.fq.gz \
    --output test_quant

# Check mapping rate
cat test_quant/logs/salmon_quant.log | grep "Mapping rate"

```

---

## Library Type Detection Fails

### Symptoms:

```
{
  "expected_format": "IU",
  "compatible_fragment_ratio": 0.45,      // Low!
  ...
}
```

### Causes & solutions:

1. Truly unstranded library:
    - OK if protocol is unstranded
    - Check library prep method
  2. Mixed orientation:
    - Sample contamination or mixing
    - Check sample sheet
  3. Very low depth:
    - Not enough reads for confident detection
    - Use --libType ISR (or appropriate type) manually
  4. Poor quality:
    - High error rate confuses detection
    - Check FastQC results
- 

## Out of Memory

Error: "std::bad\_alloc" or "Cannot allocate memory"

### Solutions:

1. Increase RAM allocation:

- Allocate 16GB minimum
- 32GB for very large transcriptomes

2. Reduce index size (rebuild with --sparse):

```
salmon index \  
    --sparse \  
        # Sparse index (less memory) \  
    --transcripts gentrome.fa \  
    --index salmon_index_sparse
```

3. Use --reduceGCMemory:

```
salmon quant \  
    --reduceGCMemory \  
        # Reduce garbage collection overhead \  
    --index salmon_index \  
    ...
```

---

### Very Low Fragment Count

**Symptoms:** Most transcripts have <10 fragments assigned

**Causes:**

1. Very low sequencing depth:

- Check total read count
- May need more sequencing

2. Wrong index:

- Using transcriptome index when should use genome
- Or vice versa

3. Severe degradation:

- Fragments too short
  - Check fragment length distribution
- 

### Inconsistent Results Across Samples

**Check:**

1. Library types match:

```
grep "expected_format" */lib_format_counts.json
```

2. Mapping rates similar:

```
grep "Mapping rate" */logs/salmon_quant.log
```

3. Fragment lengths similar:

```
# Compare fld.gz across samples
```

4. Same index used:

```
grep "index_seq_hash" */meta_info.json | sort -u
```

---

## Best Practices

### Standard RNA-seq Quantification

```
salmon quant \
  --validateMappings \
  --gcBias \
  --seqBias \
  --posBias \
  --libType A \
  --index salmon_index \
  --mates1 R1.fastq.gz \
  --mates2 R2.fastq.gz \
  --threads 4 \
  --output sample_quant
```

---

### Quick QC Run (Speed Priority)

```
salmon quant \
  --libType A \
  --index salmon_index \
  --mates1 R1.fastq.gz \
  --mates2 R2.fastq.gz \
  --threads 8 \
  --output sample_quant
# No bias corrections + 2x faster
```

---

### Publication-Quality (Accuracy Priority)

```
salmon quant \
  --validateMappings \
  --gcBias \
  --seqBias \
  --posBias \
  --numBootstraps 30 \
  --libType A
```

```
--index salmon_index \
--mates1 R1.fastq.gz \
--mates2 R2.fastq.gz \
--threads 8 \
--output sample_quant
```

---

## Related Documentation

- **Salmon Index:** [docs/salmon\\_index.md](#) - Building the index
  - **tximport:** <https://bioconductor.org/packages/tximport/> - Importing to R
  - **DESeq2:** [docs/deseq2.md](#) - Differential expression
  - **Salmon Manual:** <https://salmon.readthedocs.io/>
- 

**Document Version:** 1.0

**Last Updated:** January 2026

**Salmon Version:** 1.10.0+

**Applicable to:** Bulk RNA-seq, scRNA-seq (alevin mode), metatranscriptomics